

Properties of CCD cameras

1. Introduction

In this lab, you will measure characteristic properties of the observing equipment used at Mt. Stony Brook. You will gather sets of calibration data for both imaging and spectroscopic observations, and develop the strategy of how to calibrate science data. *Your data is always only as good as your calibration!*

1.1. CCD cameras

Charge-coupled devices (CCDs) are the sensors of choice for optical, infrared, and X-ray astronomy. At its heart, a CCD detector converts incoming photons to electrons. A big advantage of CCDs is that the response is nearly linear over most of the dynamic range, i.e. the ratio of electrons to incoming photons is nearly constant:

$$N_{\text{electrons}} \propto \text{flux} \times \Delta t \quad , \quad (1)$$

where Δt is the exposure time. (In comparison, photographic plates have an approximately logarithmic response.) These electrons are read out and digitized into “counts” using an analog-to-digital converter (ADC). To infer the flux of photons from the observed object from the measured counts requires a series of calibration steps. Most of these are based on additional calibration data which you need to acquire. Some of the calibration data can be taken during the day-time, a few types of calibration data have to be taken at night-time, under the same conditions as your science data.

The following properties of CCD cameras are important for understanding the calibration steps:

Gain:

The *gain*, which is set by the electronics of the camera, expresses the conversion of electrons into counts by the ADC:

$$\text{gain} = \frac{N_{\text{electrons}}}{N_{\text{counts}}} \quad . \quad (2)$$

Its unit is electrons per count. Typical values for the gain are $\sim 2 - 3$.

Quantum efficiency:

Ideally, every photon striking the CCD would free an electron. However, this is not the case, and the *quantum efficiency* (QE) of a CCD expresses the ratio of electrons to photons. QE is a function of wavelength. Modern CCDs reach QE < 90% in the optical. However, other components of the equipment further reduce the number of photons that reach the CCD: the reflectivity of the mirrors, transmission of lenses and filters, as well as the transmission of the atmosphere. The product of all of these is the total sensitivity, which again is a function of wavelength.

We will not measure QE or total sensitivity in this lab, but we will investigate the variation of the total sensitivity as function of position on the CCD. The total sensitivity can vary between pixels because of intrinsic QE variations, as well as varying fractions of incoming photons: e.g. because of vignetting (the further away from the optical axis, the less light is received), or dust grains blocking the light.

Bias level:

Every true CCD camera has an electronical offset level applied to it, the so-called *bias*. Hence, even if there is no signal, the counts measured in a bias frame (an exposure of 0 seconds duration) are greater than 0.

Dark current:

In cameras that are not cooled to $\lesssim -100^\circ \text{C}$, electrons can spontaneously tunnel into the conduction band even when there is no incoming light, mimicking a signal. The number of electrons is proportional to the time over which the data is integrated, i.e. the exposure time. The dark current is measured in units of electrons/pixel/second.

Dead pixels:

Pixels without a response, i.e. even with large incoming photon flux, no electrons are transferred into the conduction band.

Hot pixels:

Pixels with large dark current. These pixels will saturate even for short exposure times.

For dead pixels and hot pixels, the linear relation between counts and photons is broken. It is not possible to recover information from these pixels, and they should be removed from the analysis. This can be done with a *bad pixel map*, an image of the same dimensions as the CCD, where good pixels have a value of 1 and bad pixels have a value of 0.

Read noise:

The process of measuring the number of electrons in a pixel and converting them to counts is not perfect, it generates random noise. Read noise is measured in e^-/px .

1.2. Spectrographs

CCDs are also used as sensors in spectrographs, so very similar concepts apply. In addition, the following properties are important:

(Spectral) resolution:

The smallest wavelength difference $\Delta\lambda$ that can be distinguished at wavelength λ . This is determined by the instrumental set-up, in particular the grating, as well as the width of the entrance slit. It can be limited by the atmospheric seeing or the size of the pixels. Spectral resolution is expressed as:

$$R = \frac{\lambda}{\Delta\lambda} . \quad (3)$$

Dispersion:

The length $\Delta\lambda'$ of the spectrum over a single pixel, i.e. in units of [$\text{\AA}/\text{px}$].

1.3. Calibration data for imaging**Bias frames:**

A *bias frame* is an exposure of 0 seconds duration (and a closed shutter). They are used to measure the bias level of every pixel.

Dark frames:

Dark frames are taken with the same exposure time (and CCD temperature) as the science exposures, but with a closed shutter, i.e. no signal. They are used to measure the dark current of each pixel, and to identify hot pixels.

Flat fields:

Exposures of evenly illuminated surfaces, such as the twilight sky (*sky flats*) or illuminated parts of the dome (*dome flats*). Flat fields are used to measure, and correct for, the relative sensitivity of the CCD pixels to each other.

Standard stars:

If there are stars of known magnitude observed in the science images, we can use their measured fluxes to calibrate the fluxes of other objects in the image. If this is not the case, it is necessary to take separate exposures of *standard star fields*, ideally at the same airmass as the science exposure(s).

1.4. Calibration data for spectroscopy

Bias and/or dark frames need to be taken for spectroscopy, as well!

Flat fields:

Flat fields for spectroscopy are taken in an identical fashion as for imaging. However, in this case, the response at different wavelengths does not only reflect variations in the pixel sensitivity, but also variations in the light output from the flat-field lamps. It is common to normalize the response as function of wavelength, e.g. by fitting a spline function the observed spectrum of the dome lamps, and dividing the flat-field by it. The flat-field thus serves to correct for small-scale pixel-to-pixel variations, as well as to determine the location of the spectra on the CCD.

Arc lamp spectra:

Arc lamps are typically gas-discharge lamps of noble gases, with discrete emission line spectra. The wavelengths of these emission lines are well tabulated. By measuring the location of these lines on the CCD, one can relate position on the CCD to wavelength.

Spectrophotometric standard stars:

Calibrating the measured flux across wavelengths requires the measurement of a source with a known spectrum. Similar to the standard stars used in photometry, spectrophotometric standard stars have well tabulated spectra.

1.5. This lab

In this lab, we will measure the following properties of the CCD camera we use for imaging:

- Read noise at different temperatures
- Fraction of hot and dead pixels at different temperatures
- Bias level
- Dark current at different temperatures
- The Flat-field and variation in sensitivity

For the spectrograph, we will measure the following:

- Dispersion
- The Flat-field and variation in sensitivity

As part of the lab report, you will be asked to discuss how the different calibration steps need to be applied to the scientific data you will take in future labs.

2. Equipment

- STL1001 CCD camera (camera used for imaging)
- ST402ME CCD camera (camera used for spectroscopy)
- DADOS spectrograph with 900 lines/mm grating (“high-resolution”)
- Neon arc lamp for spectrograph
- plot of Neon emission spectrum
- laptop with CCDSoft

3. Data acquisition

3.1. Set up the imaging camera

Follow the “CCDSOft step-by-step instructions” to set up the STL-1001 camera and connect it to the control software CCDSoft. **Cool the camera to 0°C.** Note that the camera takes a while to completely cool at the center, so start cooling as soon as possible. When you’re ready to take your data, take a bias exposure to flush the CCD, and discard it. Use the “File Defaults” menu to make sure that your subsequent exposures are automatically saved. **Make sure to use 1x1 binning.**

3.2. Bias and dark frames

For bias and dark frames, the camera does not have to be mounted on the telescope. Although the camera has an internal shutter, make sure that no light gets inside the camera. Keep the camera lid on.

1. Take 1 bias frame at a CCD temperature of 0°C. Take note of the typical count levels in the images.
2. Take a series of dark frames with increasing exposure times, ranging from 10 s to ~ 2 min.

3. Take a series of 10 dark frames at a typical exposure time (~ 30 sec).
4. Take a bias frame, and a dark with your longest exposure time, at a different CCD temperature (difference of at least $\pm 10^\circ\text{C}$).

3.3. Flat fields

For taking flat fields, the camera needs to be mounted on the telescope. The type of exposure has to be set to “light”, and “auto-dark” should be selected (“auto-dark” will take a dark frame of the same exposure time and automatically subtract it from the image as a first order correction for the dark current).

4. Take a series of 10 dome flats, with counts at $\sim 30\%$ of the detector’s saturation threshold. To do so, slew dome and/or the telescope (by using the arrow keys on the hand controller) so that the telescope points at a featureless, evenly illuminated part of the dome. Take a test exposure to estimate the count rate per second, and adjust your exposure time to reach the $\sim 30\%$ target. If necessary, use the dimmer to adjust the brightness of the dome lamps.

After taking these flat-fields, disconnect and disassemble the STL-1001 camera.

3.4. Spectroscopic calibration

Assemble the spectrograph with the 20mm eyepiece viewing the spectrum. Refer to the “Spectrograph step-by-step instructions” for guidance. Attach the Neon calibration lamp to the spectrograph entrance, and observe the lines through the eyepiece. Focus the spectrograph, and change the angle of the diffraction grating so that you can see the bright yellow Neon line at 5852.48\AA (all other bright Neon lines are redder). Make sure to take turns with your lab partners.

Next, exchange the 20mm eyepiece for the ST-402 camera. Connect the camera to CCDSoft, and set up your data acquisition (temperature, auto-saving, etc.). Make sure that “auto-darks” are being taken. Take a bias frame to flush the CCD, and discard this frame. Focus the spectrograph by observing the lines from the arc lamp (use the CCDSoft Focus Tool to continuously read out images). Change the angle of the diffraction grating so that the Neon line at 5852.48\AA is located towards the “blue” end of your spectrum, i.e. your spectral range is well sampled by this line and the redder lines.

Once the spectrograph is focused, attach it to the telescope. Be careful not to touch the focuser ring!

6. Take an arc lamp spectrum by shining the arc lamp into the telescope. The dome lights should be switched off for this step.
7. Take a flat-field by turning the dome lights back on. Make sure you get enough counts above the bias level over the entire spectrum. **Keep the dome shutter and doors closed to avoid contaminating your flat-field with the Sun’s spectrum.**

4. Data analysis

The following is best done in python. The spectroscopy analysis can be done in pyraf.

4.1. Bias frames

1. Open the 0°C bias frames in python and plot a histogram¹ of the distribution of counts. Make sure to optimize the plot parameters (e.g. x-range, number of bins, logarithmic axes) in order to be able to assess the count distribution (if necessary, plot two panels with different settings). Can you identify any hot pixels? If so, define a cut that rejects them. What fraction of pixels gets rejected?
2. The standard deviation of a bias frame is a measure of the read noise in units of counts. The measured standard deviation needs to describe the bulk of the count distribution - since the distribution is not a perfect Gaussian, you will likely need to limit the count values from which you calculate the mean and the standard deviation. Determine this range, the mean and standard deviation calculated from this range, and overplot the resulting Gaussian on the count histogram. Look up the gain of the CCD camera in its header value, and convert your measurement of the read noise into units of electrons. Is it consistent with the manufacturer's description?

4.2. Dark frames

1. Make a **median** combine of the series of 10 dark frames with the same exposure time. This is a *master dark frame*. Identify the “hot” pixels in the master dark frame. Are these pixels “hot” in all the individual images with this exposure time? Are there pixels that are hot (or warm) in some individual images, but not others? What could these be?
2. For each dark frame in the time series, measure the mean and standard deviation. Make sure to limit the range of count values so that the measurements are a good description of the bulk of the count values (similar to what you did for the bias frame).
 - For each frame, determine the *statistical* uncertainty on the mean counts by considering each pixel an independent measurement of the mean counts.
 - *Systematic* uncertainties are generally more difficult to estimate. Here, you can estimate it by changing the range of count values from which you compute the mean. Which uncertainty dominates?
 - Plot the mean counts against the exposure time for your dark frames taken at 0°C. Make sure to include error bars that indicate the uncertainty. Perform a linear regression.
 - Make a second plot that shows the residuals of the data points (with uncertainties) relative to the linear regression. Is the best-fit line a good description of the data?
 - What is the dark current in electrons per pixel per second? Compare it to the manufacturer's specification.
 - Estimate the dark current for the other temperature at which you took a bias and dark frame. Explain why or why not it changes.

4.3. Imaging flat-fields

1. Take the median of your 10 flat-fields to reduce noise from photon-counting. Normalize it to its typical count level. This is your *master flat-field*.

¹ Note: remember that the count values are integers - make sure that your histogram bins make sense, e.g. a bin for values of 1.2-1.8 would not register any counts.

2. Open the master flat-field in ds9. Note how some parts of the image receive less light than others. Identify regions of particularly low counts rates and quantify what fraction of light (compared to the brightest part of the flat-field) they receive. (To help quantify this number, you can use the “Pixel Table” in ds9 found under the “Analysis” tab, and/or a “Projection” region. To place a region, you first have to select “Region” under the “Edit” menu.)
3. If you placed the same star first into the center of the image, and then close to one the corners, how would its “observed” magnitude change?
4. Plot a histogram of the counts in your master flat-field. Can you identify any dead pixels?
5. If you forgot to take flat-fields on the night of your observations, can you re-take them later?

4.4. Spectroscopic calibration

1. The spectrograph has 3 slits, with widths 25, 35 and 50 μm . Open the spectroscopic flat-field in ds9. You should see 3 spectra - which of them corresponds to the 50 μm slit?
2. In the following, you will work with the spectrum taken with the 50 μm slit. Cut out the images to keep only the data from your target slit. “Collapse” the flat-field from 2d to 1d by taking the median value of each column. Make a plot of the 1d flat-field values against the pixel positions along the dispersion axis.
3. Recall that the variation in flat-field counts with wavelength can be due to either a change in sensitivity, or the intrinsic spectrum of the lamps used to take the flat-field. Since we cannot tell the difference at this point, we will first “normalize” the flat-field along the dispersion axis to vary around 1. Variations between neighboring pixels thus reflect true differences in sensitivity, but we do not draw conclusions about variations between pixels on large pixels (which may be dominated by the lamp spectrum). To normalize the flat-field, first fit a low-order polynomial to the observed 1d flat-field spectrum². Divide the 2d flat-field by the fit - this is your normalized flat-field. Open it in ds9 and discuss how it changed compared to the original.
4. Cut out the arc lamp spectrum of the 50 μm slit (make sure that the cut-out is exactly the same as for the flat-field). Apply the normalized flat field to the arc lamp spectrum.
5. Derive the wavelength calibration from the arc lamp spectrum. To do so, identify the emission lines that you see and make a table consisting of the pixel positions and wavelengths of the lines. Plot wavelength vs. pixel position and find the best-fit line (or polynomial) to compute wavelength from pixel position - this is the wavelength calibration.
6. Plot the calibrated spectrum, labelling a few of the Neon emission lines.
7. What is the length of the spectrum (in \AA) that is covered by the spectrograph? What is the dispersion of the spectrograph?

² Note that the “wiggles” with a period of ~ 100 pixels are variations in the spectrograph response. You can choose higher-order polynomials to fit these, but your data analysis will be more straightforward you don’t.

4.5. Calibration strategies

Discuss how to apply the various calibration files. Which calibration exposures have to be taken with the same telescope + instrument set-up as the observations; which ones could you take on the next day, if necessary?

4.6. Make a pretty image

Take the image of the astronomical object of your choosing and reduce the image through the calibration steps. Make a 'pretty' image that is as good as possible exempt from artifacts. Do the calibration for all three photometric bands and produce an rgb image. Discuss imperfections in the reductions and lessons you learned to improve (if necessary) the calibration data or reduction procedure for the next labs.

5. Lab report

Prepare a `jupyter` notebook that documents your entire analysis for the lab. Make sure to explain your steps and conclusions - imagine writing a tutorial for another astronomy student who is not taking the class. Use "markdown" boxes (which can parse $\text{L}^{\text{A}}\text{T}_{\text{E}}\text{X}$; note that you can also include figures produced outside of the notebook, e.g. with `ds9`) for these explanations - this is what we will read (we will only look at your code if we think you did something wrong). Make sure that the report is logical (i.e. each section should have a short introduction, then code with results and plots, then conclusion), and make sure the section numbering follows this manual. Once your notebook is finished, make sure to restart it and re-run all cells. Then save the notebook in pdf format, e.g. through the print menu.

The timeline for the lab report, and intermediate check-ins is the following:

- +1 week: Complete Sect. 4.1 and Sect. 4.2 and check your results, and your current notebook, with a TA / the instructor.
- +2 weeks: Complete Sect. 4.3 and Sect. 4.4.1-4.4.4 and check your results with a TA / the instructor.
- +3 weeks: Complete the remainder of Sect. 4.4, and Sect. 4.5, and make a pretty image (Sect. 4.6) .
- +4 weeks: Hand in your lab report.